

Introduction to Computer Networks

Polly Huang

EE NTU

Unix Network Programming

The socket
struct and data handling
System calls

Based on **Beej's Guide to Network Programming**

Quiz Time!

The Unix Socket

- A file descriptor really
- The Unix fact
 - When Unix programs do any sort of I/O, they do it by reading or writing to a file descriptor
 - A file descriptor is simply an integer associated with an open file

A File Descriptor

- A file in Unix can be
 - A network connection
 - A FIFO queue
 - A pipe
 - A terminal
 - A real on-the-disk file
 - Or just about anything else

Jeez, *everything* in Unix is a file!

Well, we know how to handle files!

- In theory
 - The `read()` and `write()` calls allows to communicate through a socket
- In practice
 - The `send()` and `recv()` offer much greater control over your data transmission

The structs

- **int**
 - For the file descriptor
- **struct sockaddr**
 - Space holder for “types” of addresses
- **struct sockaddr_in**
 - Specific for the “Internet” type
 - **_in** for Internet
- **struct in_addr**
 - 4 byte IP address

struct sockaddr

```
struct sockaddr {  
    unsigned short sa_family;  
        // address family, AF_XXX  
    char sa_data[14];  
        // 14 bytes of protocol address  
};
```

struct sockaddr_in

```
struct sockaddr_in {  
    short int sin_family;  
        // Address family unsigned, AF_INET  
    short int sin_port;  
        // Port number, in network byte order  
    struct in_addr sin_addr;  
        // Internet address, in network byte order  
    unsigned char sin_zero[8];  
        // Same size as struct sockaddr  
};
```

Struct in_addr

```
struct in_addr {  
    // Internet address (a structure for historical reasons)  
    unsigned long s_addr;  
    // that's a 32-bit long, or 4 bytes  
};
```

Reference

- Let *ina* be of type `struct sockaddr_in`
- *ina.sin_addr.s_addr* references the 4-byte IP address in network byte order

Types of Byte Ordering

- Network Byte Order
 - Most significant byte first
 - Need conversion from the app program to the network
- Host Byte Order
 - Least significant byte first
 - Usually no need in app program
 - But need conversion if data coming from the network

Functions to Convert

- `htons()`
 - Host to Network Short
- `htonl()`
 - Host to Network Long
- `ntohs()`
 - Network to Host Short
- `ntohl()`
 - Network to Host Long

Storing the IP address

```
ina.sin_addr.s_addr = inet_addr("10.12.110.57");
```

- Returns “-1” on error
- For unsigned short it’s 255.255.255.255
- A broadcast address

A Cleaner Interface

- `#include <sys/socket.h>`
- `#include <netinet/in.h>`
- `#include <arpa/inet.h>`
- `int inet_aton(const char *cp, struct in_addr *inp);`

An Example

```
struct sockaddr_in my_addr;  
my_addr.sin_family = AF_INET;  
    // host byte order  
my_addr.sin_port = htons(MYPORT);  
    // short, network byte order  
inet_aton("10.12.110.57", &(my_addr.sin_addr));  
memset(&(my_addr.sin_zero), '\0', 8);  
    // zero the rest of the struct
```

Things to Note

- `inet_addr()` and `inet_aton()` both convert IP addresses into the network byte order
- Not all platforms implement `inet_aton()`

Get the IP Address Back

- `printf("%s", inet_ntoa(ina.sin_addr));`
- `inet_ntoa()` returns a pointer to a `char*`
- And...

Use strcpy()

```
char *a1, *a2; . .  
a1 = inet_ntoa(ina1.sin_addr);  
    // this is 192.168.4.14  
a2 = inet_ntoa(ina2.sin_addr);  
    // this is 10.12.110.57  
printf("address 1: %s\n",a1);  
printf("address 2: %s\n",a2);
```

This program will print:

```
address 1: 10.12.110.57  
address 2: 10.12.110.57
```

Quiz Time!

System Calls

socket()

Creating the File Descriptor

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol);
```

domain: AF_INET

type: SOCK_STREAM or SOCK_DGRAM

protocol: 0 or getprotobyname()

bind()

Associating Port with the FD

- `#include <sys/types.h>`
- `#include <sys/socket.h>`
- `int bind(int sockfd, struct sockaddr *my_addr, int addrlen);`

Example (Typical Server)

```
#include <string.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

#define MYPORT 3490
main() {
    int sockfd;
    struct sockaddr_in my_addr;
    sockfd = socket(AF_INET, SOCK_STREAM, 0); // do some error checking!
    my_addr.sin_family = AF_INET; // host byte order
    my_addr.sin_port = htons(MYPORT); // short, network byte order
    my_addr.sin_addr.s_addr = inet_addr("10.12.110.57");
    memset(&(my_addr.sin_zero), '\0', 8); // zero the rest of the struct

    // don't forget your error checking for bind():
    bind(sockfd, (struct sockaddr *)&my_addr, sizeof(struct sockaddr));
    . . .
```

connect()

Making a Connection

- `#include <sys/types.h>`
- `#include <sys/socket.h>`
- `int connect(int sockfd, struct sockaddr *serv_addr, int addrlen);`

Example (Typical Client)

```
#include <string.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

#define DEST_IP "10.12.110.57"
#define DEST_PORT 23
main() {
    int sockfd;
    struct sockaddr_in dest_addr; // will hold the destination addr
    sockfd = socket(AF_INET, SOCK_STREAM, 0); // do some error checking!
    dest_addr.sin_family = AF_INET; // host byte order
    dest_addr.sin_port = htons(DEST_PORT); // short, network byte order
    dest_addr.sin_addr.s_addr = inet_addr(DEST_IP);
    memset(&(dest_addr.sin_zero), '\0', 8); // zero the rest of the struct

    // don't forget to error check the connect()!
    connect(sockfd, (struct sockaddr *)&dest_addr, sizeof(struct sockaddr));
    . . .
```

listen()

Waiting for Connection

```
#include <sys/socket.h>
```

```
int listen(int sockfd, int backlog);
```

On the server side, you see typically this:

```
socket();
```

```
bind();
```

```
listen();
```

```
/* accept() goes here */
```

accept()

Getting a Connection

```
#include <sys/socket.h>
```

```
int accept(int sockfd, void *addr, int *addrlen);
```

The Server Example

```
#include <string.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

#define MYPORT 3490 // the port users will be connecting to
#define BACKLOG 10 // how many pending connections queue will hold
main() {
    int sockfd, new_fd; // listen on sockfd, new connection on new_fd
    struct sockaddr_in my_addr; // my address information
    struct sockaddr_in their_addr; // connector's address information
    int sin_size;

    sockfd = socket(AF_INET, SOCK_STREAM, 0); // do some error checking!
    my_addr.sin_family = AF_INET; // host byte order
    my_addr.sin_port = htons(MYPORT); // short, network byte order
    my_addr.sin_addr.s_addr = INADDR_ANY; // auto-fill with my IP
    memset(&(my_addr.sin_zero), '\0', 8); // zero the rest of the struct

    // don't forget your error checking for these calls:
    bind(sockfd, (struct sockaddr *)&my_addr, sizeof(struct sockaddr));
    listen(sockfd, BACKLOG);
    sin_size = sizeof(struct sockaddr_in);
    new_fd = accept(sockfd, (struct sockaddr *)&their_addr, &sin_size);
```

send() and recv() Data Transmission

```
int send(int sockfd, const void *msg, int len, int flags);  
int recv(int sockfd, void *buf, int len, unsigned int flags);
```

Example

```
char *msg = "Hello World!";  
int len, bytes_sent;  
..  
len = strlen(msg);  
bytes_sent = send(sockfd, msg, len, 0);  
...
```


Quiz Time!

`sendto()` and `recvfrom()`

Transmission the Datagram Style

```
int sendto(int sockfd, const void *msg, int len, unsigned  
int flags, const struct sockaddr *to, int tolen);
```

```
int recvfrom(int sockfd, void *buf, int len, unsigned int  
flags, struct sockaddr *from, int *fromlen);
```

Or if transmitting over **TCP socket**,
one can simply use `send()` and `recv()`.

`close()` and `shutdown()`

Closing the Communication

```
close(sockfd);
```

```
int shutdown(int sockfd, int how);
```

- 0 -- Further receives are disallowed
- 1 -- Further sends are disallowed
- 2 -- Further sends and receives are disallowed (like `close()`)

Reference

- **Beej's Guide to Network Programming**
 - <https://beej.us/guide/bgnet/>
- Additional system calls
- TCP stream client, server example
- UDP datagram example